


# Multilayer Percetrons



*A dataset*

<i>Fields</i>	<i>class</i>
2.5 3.8 2.0	0
4.9 4.5 4.3	0
7.5 3.9 2.8	1
3.6 1.2 1.3	0
etc ...	



## *Training the neural network*

***Fields***                      ***class***


2.5   3.8   2.0                0

4.9   4.5   4.3                0

7.5   3.8   2.8                1

3.6   1.2   1.3                0

etc ...





## *Training the neural network*

***Fields***                      ***class***

2.5   3.8   2.0              0


4.9   4.5   4.3              0

7.5   3.8   2.8              1

3.6   1.2   1.3              0

etc ...

**Initialise with random weights**







*Training the neural network*

**Fields** **class**

2.5	3.8	2.0	0
-----	-----	-----	---


4.9	4.5	4.3	0
-----	-----	-----	---

7.5	3.8	2.8	1
-----	-----	-----	---

3.6	1.2	1.3	0
-----	-----	-----	---

etc ...

**Input the training pattern**





*Training the neural network*

***Fields***                      ***class***

2.5	3.8	2.0	0
-----	-----	-----	---


4.9	4.5	4.3	0
-----	-----	-----	---

7.5	3.8	2.8	1
-----	-----	-----	---

3.6	1.2	1.3	0
-----	-----	-----	---

etc ...

**Feed it through to get output**





## *Training the neural network*

***Fields***                      ***class***

2.5	3.8	2.0	0
-----	-----	-----	---


4.9	4.5	4.3	0
-----	-----	-----	---

7.5	3.8	2.8	1
-----	-----	-----	---

3.6	1.2	1.3	0
-----	-----	-----	---

etc ...

**Compare with target output**








## Training the neural network

*Fields* *class*

2.5	3.8	2.0	0
4.9	4.5	4.3	0
7.5	3.8	2.8	1
3.6	1.2	1.3	0
etc ...			

**Adjust weights based on error**







## *Training the neural network*

**Fields**                      **class**

2.5 3.8 2.0                0


4.9 4.5 4.3                0

7.5 3.8 2.8                1

3.6 1.2 1.3                0

etc ...

**Present a training pattern**



*Training the neural network*

***Fields***                      ***class***

2.5 3.8 2.0                      0


4.9 4.5 4.3                      0

7.5 3.8 2.8                      1

3.6 1.2 1.3                      0

etc ...

**Feed it through to get output**





*Training the neural network*

**Fields**                      **class**

2.5 3.8 2.0              0


4.9 4.5 4.3              0

7.5 3.8 2.8              1

3.6 1.2 1.3              0

etc ...

**Compare with target output**







## Training the neural network

**Fields** **class**

2.5 3.8 2.0 0


4.9 4.5 4.3 0

7.5 3.8 2.8 1

3.6 1.2 1.3 0

etc ...

**Adjust weights based on error**







## *Training the neural network*

***Fields***                      ***class***

2.5   3.8   2.0            0


4.9   4.5   4.3            0

7.5   3.8   2.8            1

3.6   1.2   1.3            0

etc ...

**And so on ....**




Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments

*Algorithms for weight adjustment are designed to make changes that will reduce the error*

# The decision boundary perspective...


**Initial random weights**





# The decision boundary perspective...

**Present a training instance / adjust the weights**








# The decision boundary perspective...

**Present a training instance / adjust the weights**








# The decision boundary perspective...


**Present a training instance / adjust the weights**






# The decision boundary perspective...

**Present a training instance / adjust the weights**



# The decision boundary perspective...


## Eventually ...





# Learning Algo.,

- Weight-learning algorithms for NNs
- They work by making thousands and thousands of tiny adjustments, each making the network do better at the most recent pattern, but perhaps a little worse on many others
- But, by dumb luck, eventually this tends to be good enough to  
learn effective classifiers for many real applications







# Some other points


## **Detail** of a standard NN weight learning algorithm

If  $f(x)$  is non-linear, a network with 1 hidden layer can, in theory, learn perfectly any classification problem. A set of weights exists that can produce the targets from the inputs. The problem is finding them.



# Some other ‘by the way’ points


If  $f(x)$  is linear, the NN can **only** draw straight decision boundaries (even if there are many layers of units)





# Some other ‘by the way’ points

NNs use nonlinear  $f(x)$  so they can draw complex boundaries, but keep the data unchanged










# NN and Back Propagation Algorithm



- Single layer nets have limited representation power (linear Separability problem). Multi layer nets (or nets with non-linear hidden units) may overcome linear inseparability problem.
- Every boolean function can be represented by a network with a single hidden layer
- Every bounded continuous function can be approximated with arbitrary small error, by network with one hidden layer.
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers.

# Multilayer Perceptrons Architecture



# A solution for the XOR problem

$x_1$	$x_2$	$x_1 \text{ XOR } x_2$
-1	-1	-1
-1	1	1
1	-1	1
1	1	-1




$$\varphi(v) = \begin{cases} 1 & \text{if } v > 0 \\ -1 & \text{if } v \leq 0 \end{cases}$$

$\varphi$  is the sign function.

# NEURON MODEL

- Sigmoidal Function



$$\varphi(v_j) = \frac{1}{1 + e^{-av_j}}$$


$$v_j = \sum_{i=0, \dots, m} w_{ji} y_i$$

- $v_j$  induced field of neuron  $j$
- Most common form of activation function
- $a \rightarrow \infty \Rightarrow \varphi \rightarrow$  threshold function
- Differentiable*



# LEARNING ALGORITHM

- Back-propagation algorithm



- It adjusts the weights of the NN in order to minimize the average squared error.

# Average Squared Error

- Error signal of output neuron  $j$  at presentation of  $n$ -th training example:

$$e_j(n) = d_j(n) - y_j(n)$$

- Total error at time  $n$ :

$$E(n) = \frac{1}{2} \sum_{j \in C} e_j^2(n)$$

- Average squared error:

- Measure of learning performance:

$$E_{AV} = \frac{1}{N} \sum_{n=1}^N E(n)$$

$C$ : Set of neurons in output layer

$N$ : size of training set

- **Goal:** *Adjust weights of NN to minimize  $E_{AV}$*

# Notation

$e_j$  Error at output of neuron  $j$

$y_j$  Output of neuron  $j$

$v_j = \sum_{i=0, \dots, m} w_{ji} y_i$  Induced local field of neuron  $j$

# Weight Update Rule


Update rule is based on the gradient descent method  
take a step in the direction yielding the maximum  
decrease of E

$$\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}}$$

Step in direction opposite to the gradient



# Computing Model (output neuron)



Signal-flow graph highlighting the details of output neuron  $j$ .

# Definition of the Local Gradient of neuron j

$$\delta_j = -\frac{\partial E}{\partial v_j} \quad \text{Local Gradient}$$

We obtain

$$\delta_j = e_j \varphi'(v_j)$$

because

$$-\frac{\partial E}{\partial v_j} = -\frac{\partial E}{\partial e_j} \frac{\partial e_j}{\partial y_j} \frac{\partial y_j}{\partial v_j} = -e_j(-1)\varphi'(v_j)$$

# Update Rule

- We obtain

$$\Delta w_{ji} = \eta \delta_j y_i$$

because

$$\frac{\partial E}{\partial w_{ji}} = \frac{\partial E}{\partial v_j} \frac{\partial v_j}{\partial w_{ji}}$$

$$-\frac{\partial E}{\partial v_j} = \delta_j \quad \frac{\partial v_j}{\partial w_{ji}} = y_i$$

# Compute local gradient of neuron $j$

- The key factor is the calculation of  $e_j$
- There are two cases:
  - Case 1):  $j$  is a output neuron
  - Case 2):  $j$  is a hidden neuron



# Error $e_j$ of output neuron

- Case 1:  *$j$  output neuron*

$$e_j = d_j - y_j$$


Then

$$\delta_j = (d_j - y_j)\varphi'(v_j)$$

# Local gradient of hidden neuron

- Case 2: ***j hidden neuron***
- the local gradient for neuron  $j$  is recursively determined in terms of the local gradients of all neurons to which neuron  $j$  is directly connected

# Computing model (hidden neuron)



Signal-flow graph highlighting the details of output neuron  $k$  connected to hidden neuron  $j$ .

# Use the Chain Rule

$$\delta_j = -\frac{\partial E}{\partial y_j} \frac{\partial y_j}{\partial v_j} \qquad \frac{\partial y_j}{\partial v_j} = \varphi'(v_j)$$

$$E(n) = \frac{1}{2} \sum_{k \in C} e_k^2(n)$$

$$-\frac{\partial E}{\partial y_j} = -\sum_{k \in C} e_k \frac{\partial e_k}{\partial y_j} = \sum_{k \in C} e_k \left[ \frac{-\partial e_k}{\partial v_k} \right] \frac{\partial v_k}{\partial y_j}$$

from  $-\frac{\partial e_k}{\partial v_k} = \varphi'(v_k) \qquad \frac{\partial v_k}{\partial y_j} = w_{kj}$


We obtain  $-\frac{\partial E}{\partial y_j} = \sum_{k \in C} \delta_k w_{kj}$



# Local Gradient of hidden neuron $j$

Hence

$$\delta_j = \varphi'(v_j) \sum_{k \in C} \delta_k w_{kj}$$



Signal-flow  
graph of  
back-  
propagation  
error signals  
to neuron  $j$

# Delta Rule

- **Delta rule**  $\Delta \mathbf{w}_{ji} = \eta \delta_j \mathbf{y}_i$

$$\delta_j = \begin{cases} \varphi'(v_j)(d_j - y_j) & \text{IF } j \text{ output node} \\ \varphi'(v_j) \sum_{k \in C} \delta_k w_{kj} & \text{IF } j \text{ hidden node} \end{cases}$$

C: Set of neurons in the layer following the one containing  $j$

# Local Gradient of neurons

$$\varphi'(v_j) = ay_j[1 - y_j]$$

$$a > 0$$

$$\delta_j = \begin{cases} ay_j[1 - y_j] \sum \delta_k w_{kj} & \text{if } j \text{ hidden node} \\ ay_j[1 - y_j][d_j^k - y_j] & \text{if } j \text{ output node} \end{cases}$$


# Backpropagation algorithm

- Two phases of computation:
  - **Forward pass**: run the NN and compute the error for each neuron of the output layer.
  - **Backward pass**: start at the output layer, and pass the errors backwards through the network, layer by layer, by recursively computing the local gradient of each neuron.



# Summary

## Multilayer Perceptrons



Signal-flow graphical summary of back-propagation learning.  
 Top part of the graph: forward pass. Bottom part of the graph:  
 backward pass.

# Training

- **Sequential mode** (on-line, pattern or stochastic mode):
  - $(x(1), d(1))$  is presented, a sequence of forward and backward computations is performed, and the weights are updated using the **delta rule**.
  - Same for  $(x(2), d(2)), \dots, (x(N), d(N))$ .


# Training

- The learning process continues on an epoch-by-epoch basis until the stopping condition is satisfied.
- From one epoch to the next choose a **randomized** ordering for selecting examples in the training set.

# Stopping criteria

- **Sensible stopping criteria:**
  - **Average squared error change:**  
Back-prop is considered to have converged when the absolute rate of change in the average squared error per epoch is sufficiently small (in the range  $[0.1, 0.01]$ ).
  - **Generalization based criterion:**  
After each epoch the NN is tested for generalization. If the generalization performance is adequate then stop.

# Early stopping







# Generalization

- **Generalization**: NN generalizes well if the I/O mapping computed by the network is nearly correct for new data (test set).
- Factors that influence generalization:
  - the size of the training set.
  - the architecture of the NN.
  - the complexity of the problem at hand.
- **Overfitting (overtraining)**: when the NN learns too many I/O examples it may end up memorizing the training data.

# Generalization



(a) Properly fitted data (good generalization)  
(b) Overfitted data (poor generalization).

# Expressive capabilities of NN

## Boolean functions:

- Every boolean function can be represented by network with single hidden layer
- but might require exponential hidden units

## Continuous functions:

- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer
- Any function can be approximated with arbitrary accuracy by a network with two hidden layers

# Generalized Delta Rule

- If  $\eta$  small  $\Rightarrow$  Slow rate of learning  
If  $\eta$  large  $\Rightarrow$  Large changes of weights  
 $\Rightarrow$  NN can become unstable  
(oscillatory)
- Method to overcome above drawback:  
**include a momentum term in the delta rule**

$$\Delta w_{ji}(n) = \alpha \Delta w_{ji}(n-1) + \eta \delta_j(n) y_i(n)$$

*momentum constant*

*Generalized  
delta  
function*

# Generalized delta rule

- the momentum accelerates the descent in steady downhill directions.
- the momentum has a stabilizing effect in directions that oscillate in time.



# $\eta$ adaptation

Heuristics for accelerating the convergence of the back-prop algorithm through  $\eta$  adaptation:

- **Heuristic 1:** Every weight should have its own  $\eta$ .
- **Heuristic 2:** Every  $\eta$  should be allowed to vary from one iteration to the next.

# NN DESIGN

- Data representation
- Network Topology
- Network Parameters
- Training
- Validation

# Setting the parameters

- How are the weights initialised?
- How is the learning rate chosen?
- How many hidden layers and how many neurons?
- Which activation function ?
- How to preprocess the data ?
- How many examples in the training data set?

# Some heuristics (1)

- Sequential v/s Batch algorithms:
- the sequential mode (pattern by pattern) is computationally faster than the batch mode (epoch by epoch)
- Sequential also called as on line learning
- Sequential is stochastic in nature
- Tracks small changes in training data set.
- Simple to implement
- Effective solution to large scale and complex classification problems
- Cost function is instantaneous error energy

## Some heuristics (2)

- Maximization of information content: every training example presented to the back propagation algorithm must maximize the information content.
  - The use of an example that results in the largest training error.
  - The use of an example that is radically different from all those previously used.

## Some heuristics (3)

- Activation function: network learns faster with antisymmetric functions when compared to nonsymmetric functions.

$$\varphi(v) = \frac{1}{1 + e^{-av}}$$


Sigmoidal function is nonsymmetric

$$\varphi(v) = a \tanh(bv)$$


Hyperbolic tangent function is nonsymmetric



# Some heuristics (3)



(a)



(b)

Antisymmetric activation function. (b) Nonsymmetric activation function.

## Some heuristics (4)

- Target values: target values must be chosen within the range of the sigmoidal activation function.
- Otherwise, hidden neurons can be driven into saturation which slows down learning

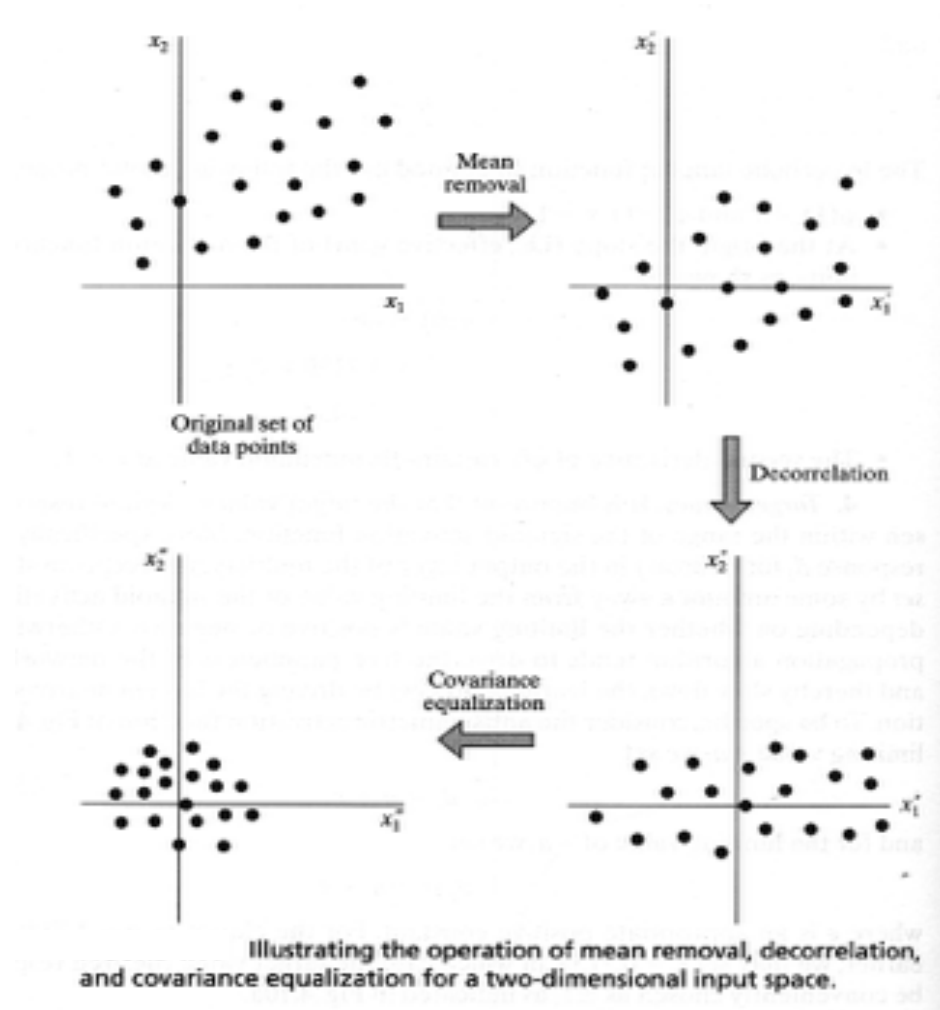
## Some heuristics (4)

- For the antisymmetric activation function it is necessary to design  $\epsilon$
- For  $a^+$ :  
$$d_j = a - \epsilon$$
- For  $-a$ :  
$$d_j = -a + \epsilon$$
- If  $a=1.7159$  we can set  $\epsilon=0.7159$   
then  $d=\pm 1$

# Some heuristics (5)

- Inputs normalisation:
  - Each input variable should be processed so that the mean value is small or close to zero or at least very small when compared to the standard deviation.
  - Input variables should be uncorrelated.
  - Decorrelated input variables should be scaled so their covariances are approximately equal.

# Some heuristics (5)



# Some heuristics (6)

- Initialization of weights:
  - If synaptic weights are assigned large initial values neurons are driven into saturation. Local gradients become small so learning rate becomes small.
  - If synaptic weights are assigned small initial values algorithms operate around the origin. For the hyperbolic activation function the origin is a saddle point.

# Some heuristics (7)

- Learning rate:
  - The right value of  $\eta$  depends on the application. Values between 0.1 and 0.9 have been used in many applications.
  - Other heuristics adapt  $\eta$  during the training as described in previous slides.



# Some heuristics (8)

- How many layers and neurons
  - The number of layers and of neurons **depend on the specific task**. In practice this issue is solved by trial and error.
  - Two types of adaptive algorithms can be used:
    - **start from a large network and successively remove some neurons and links until network performance degrades.**
    - **begin with a small network and introduce new neurons until performance is satisfactory.**

# Some heuristics (9)

- How many training data ?
  - Rule of thumb: the number of training examples should be at least five to ten times the number of weights of the network.